

Lecture 2: XSEDE Jetstream, MapReduce Programming Model, and Problem Solving

COSC 526: Introduction to Data Mining
Spring 2020



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

Instructors:

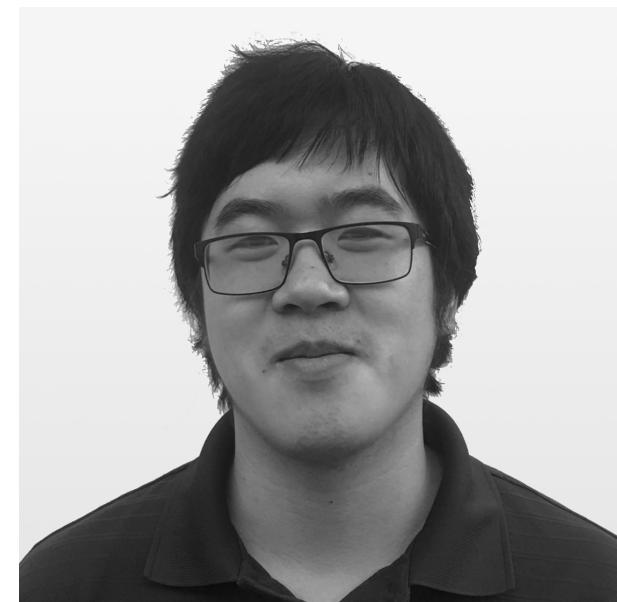


Michela Taufer



Danny Rorabaugh

GRA:



Nigel Tan

Topics covered today

- The XSEDE Jetstream cloud
 - Create your XSEDE account
- The MapReduce Programming model
 - Capture high level concepts
- Problem solving
 - Learn how different solutions can have different execution times

Jetstream

-

a brief introduction



NSF Funding Areas in HPC

Traditionally concentrated on enabling petascale capability

- Blue Waters – 13.3 petaflops, 2012
- Stampede – 9.6 petaflops, 2013
- Comet – ~2.0 petaflops, 2014

Has funded research into building clouds and computer science

- CloudLab
- Chameleon

Now funding clouds to do research

- Bridges (Hybrid system)
- Jetstream



What is Jetstream and why does it exist?

- NSF's first production cloud facility
- Part of the NSF eXtreme Digital (XD) program
- Provides on-demand *interactive* computing and analysis
- Enables *configurable* environments and architectures
- User-friendly, widely accessible cloud environment
- User-selectable library of preconfigured virtual machines



From slides by Jeremy Fischer
jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

HPC vs Cloud

Adapting to a different environment:

- No reservations, no queueing
- More interactive use and less/no batch queuing
- What? No parallel filesystem?!?
- Being your own admin – hey, we have root!
- You really can have almost any (linux) software you want**
- Constantly getting new features (<https://www.openstack.org/software/project-navigator/>)

** Here there be dragons...



From slides by Jeremy Fischer
jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

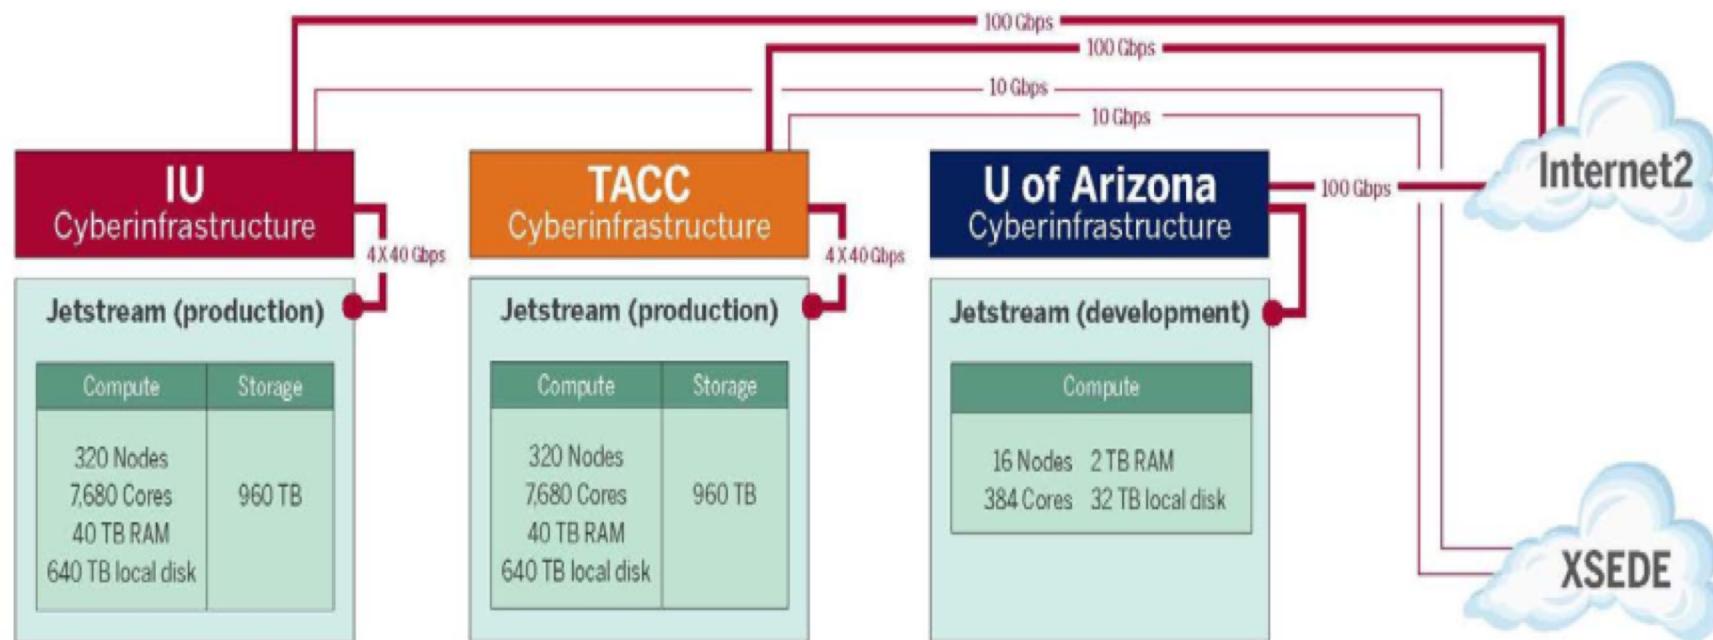
Who uses Jetstream?

- The researcher needing a handful of cores (1 to 44/vCPU)
- Software creators and researchers needing to create their own customized virtual machines and workflows
- Science gateway creators using Jetstream as either the frontend or processor for scientific jobs
- STEM Educators teaching on a variety of subjects



From slides by Jeremy Fischer
jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

Jetstream System Overview



From slides by Jeremy Fischer
jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

Hardware and Instance "Flavors"

VM Host Configuration

- Dual Intel E-2680v3 "Haswell"
- 24 physical cores/node @ 2.5 GHz (Hyperthreading on)
- 128 GB RAM
- Dual 1 TB local disks
- 10GB dual uplink NIC
- Running KVM Hypervisor

Flavor	vCPUs	RAM	Storage	Per Node
m.tiny	1	2	8	46
m.small	2	4	20	23
m.medium	6	16	60	7
m.large	10	30	120	4
m.xlarge	24	60	240	2
m.xxlarge	44	120	480	1

- Short-term storage comes as part of launched instance
- Long-term storage is XSEDE-allocated
- Implemented on backend as OpenStack Volumes
- Each user gets 10 volumes up to 500GB total storage
- Piloting object storage as well after recent update



Requesting access to Jetstream

- You can request startup allocations anytime. (Startups are simple!)
- You can request allocations for educational use anytime.



From slides by Jeremy Fischer
jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

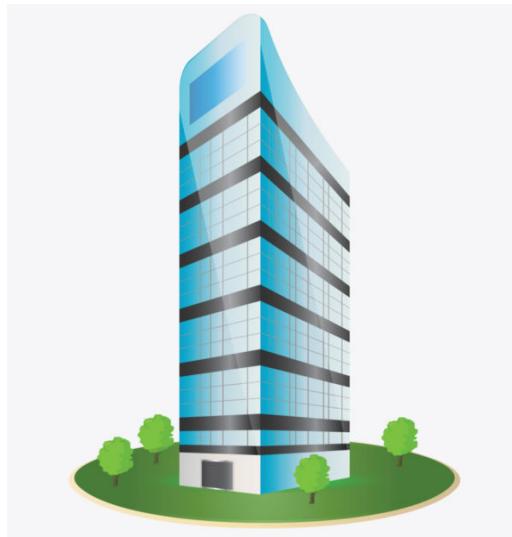
Create your XSEDE Jetstream account

- Create your account by following the directions in the file:
JetstreamGuide.ipynb
- Enter your XSEDE Jetstream user name here:
<https://tinyurl.com/wwfcysm>
- Your account will be active in a couple of days

The Canonical MapReduce Programming Model

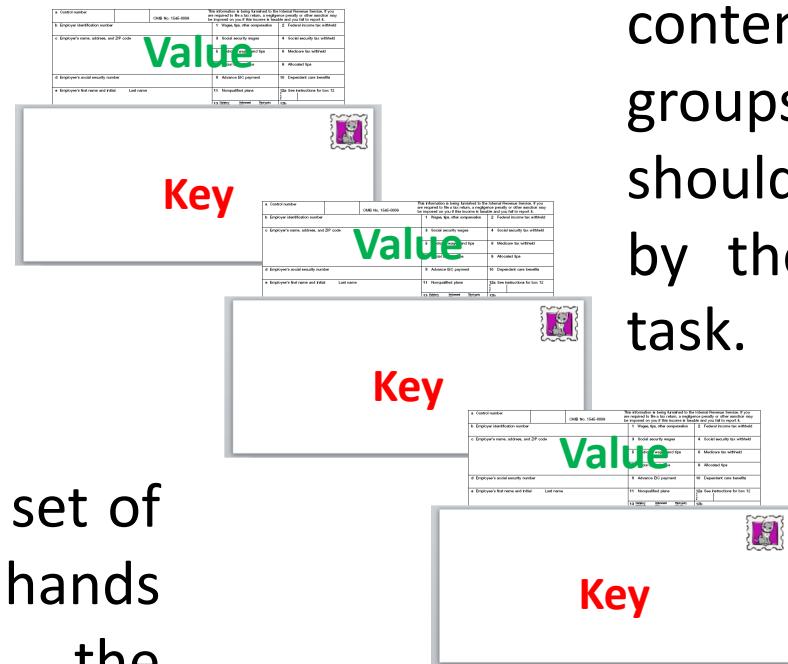


The MapReduce Workflow



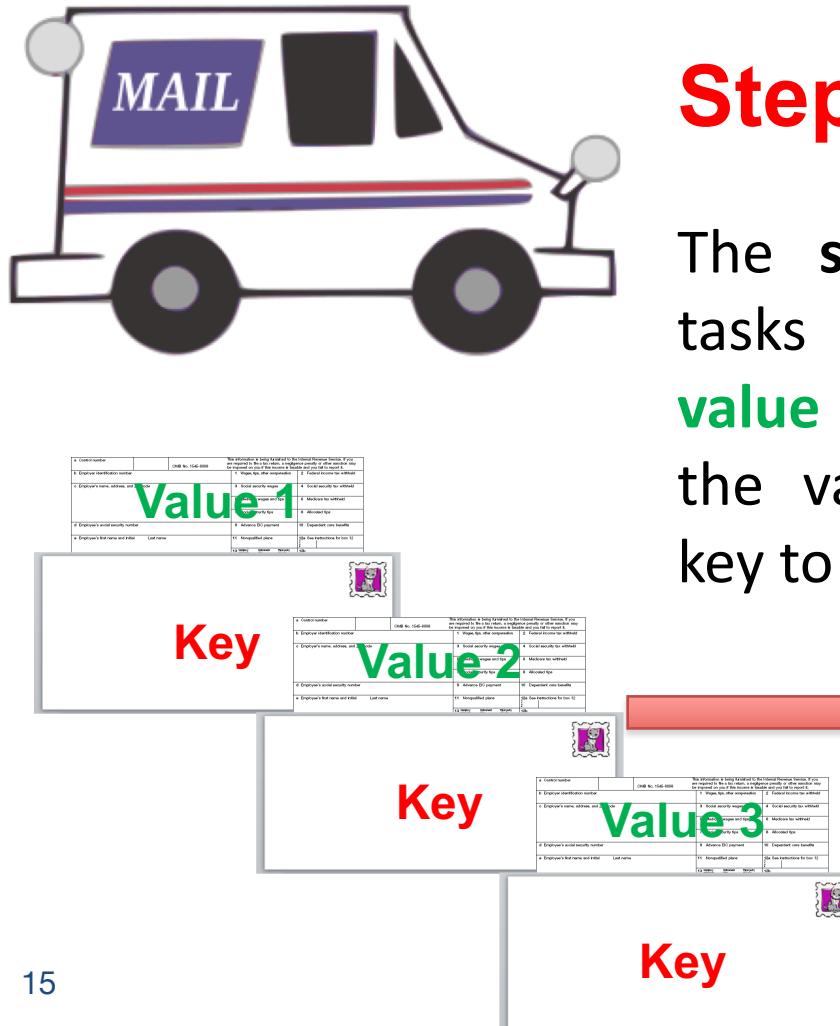
Step 1: Map

The **map** task creates a set of **Key-Value** pairs and hands them off to the **sorting/shuffling** task.



The **Value** is the content, the **Key** groups content that should be processed by the same reduce task.

The MapReduce Workflow



Step 2: Sort/Shuffle

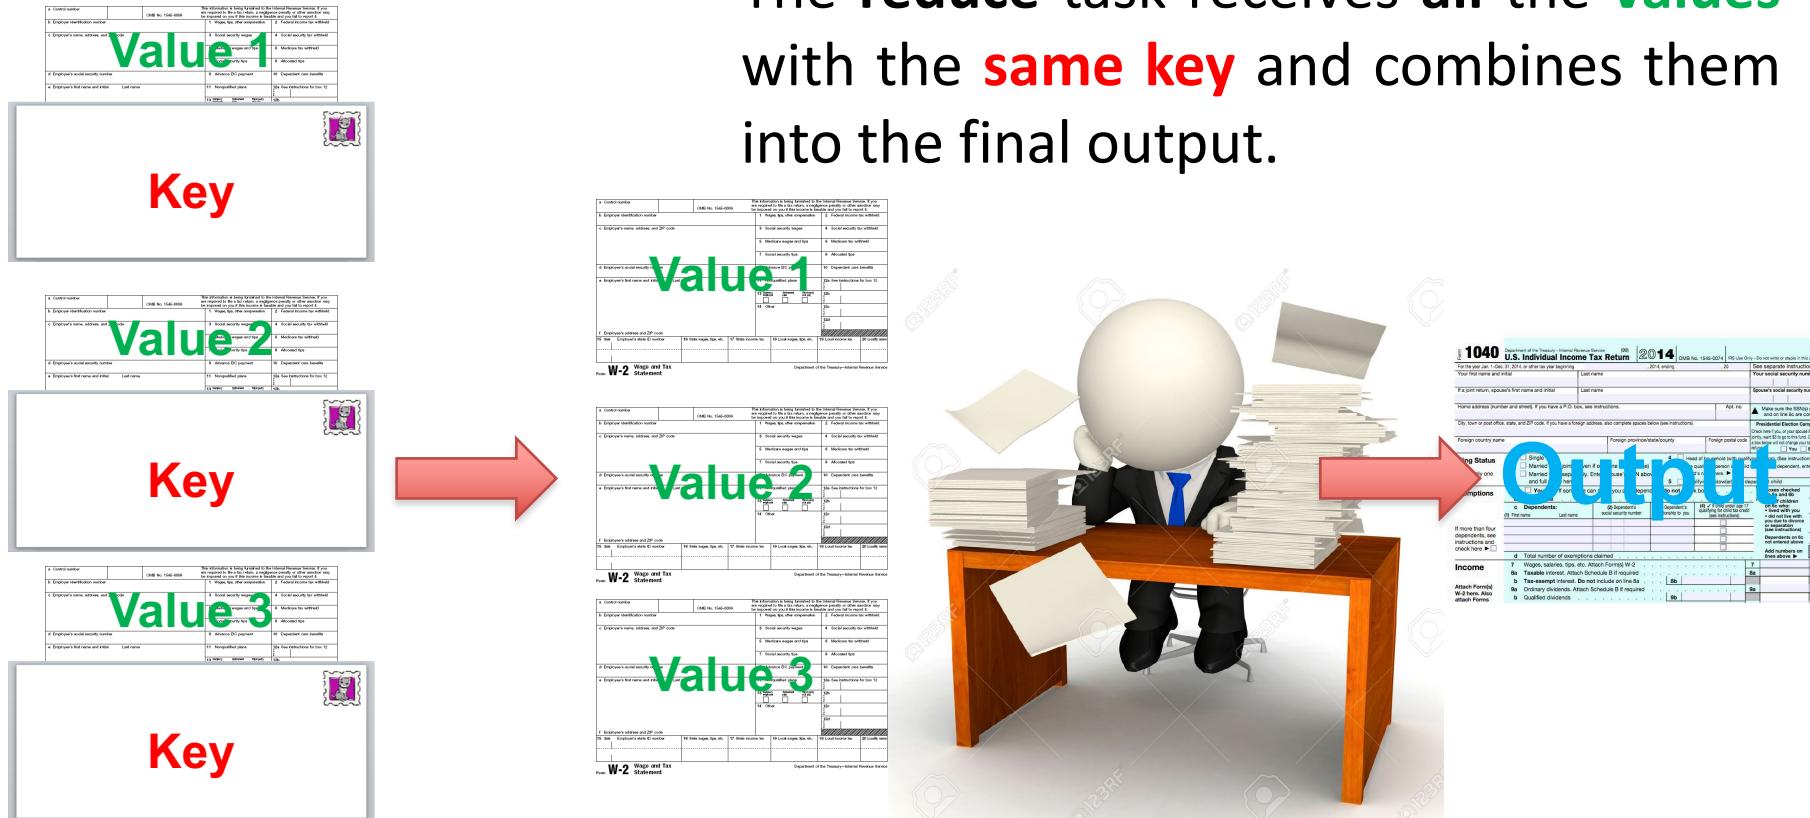
The **sorting** and **shuffling** tasks collect all the **key-value** pairs and deliver all the values with the same key to a **reduce** task.



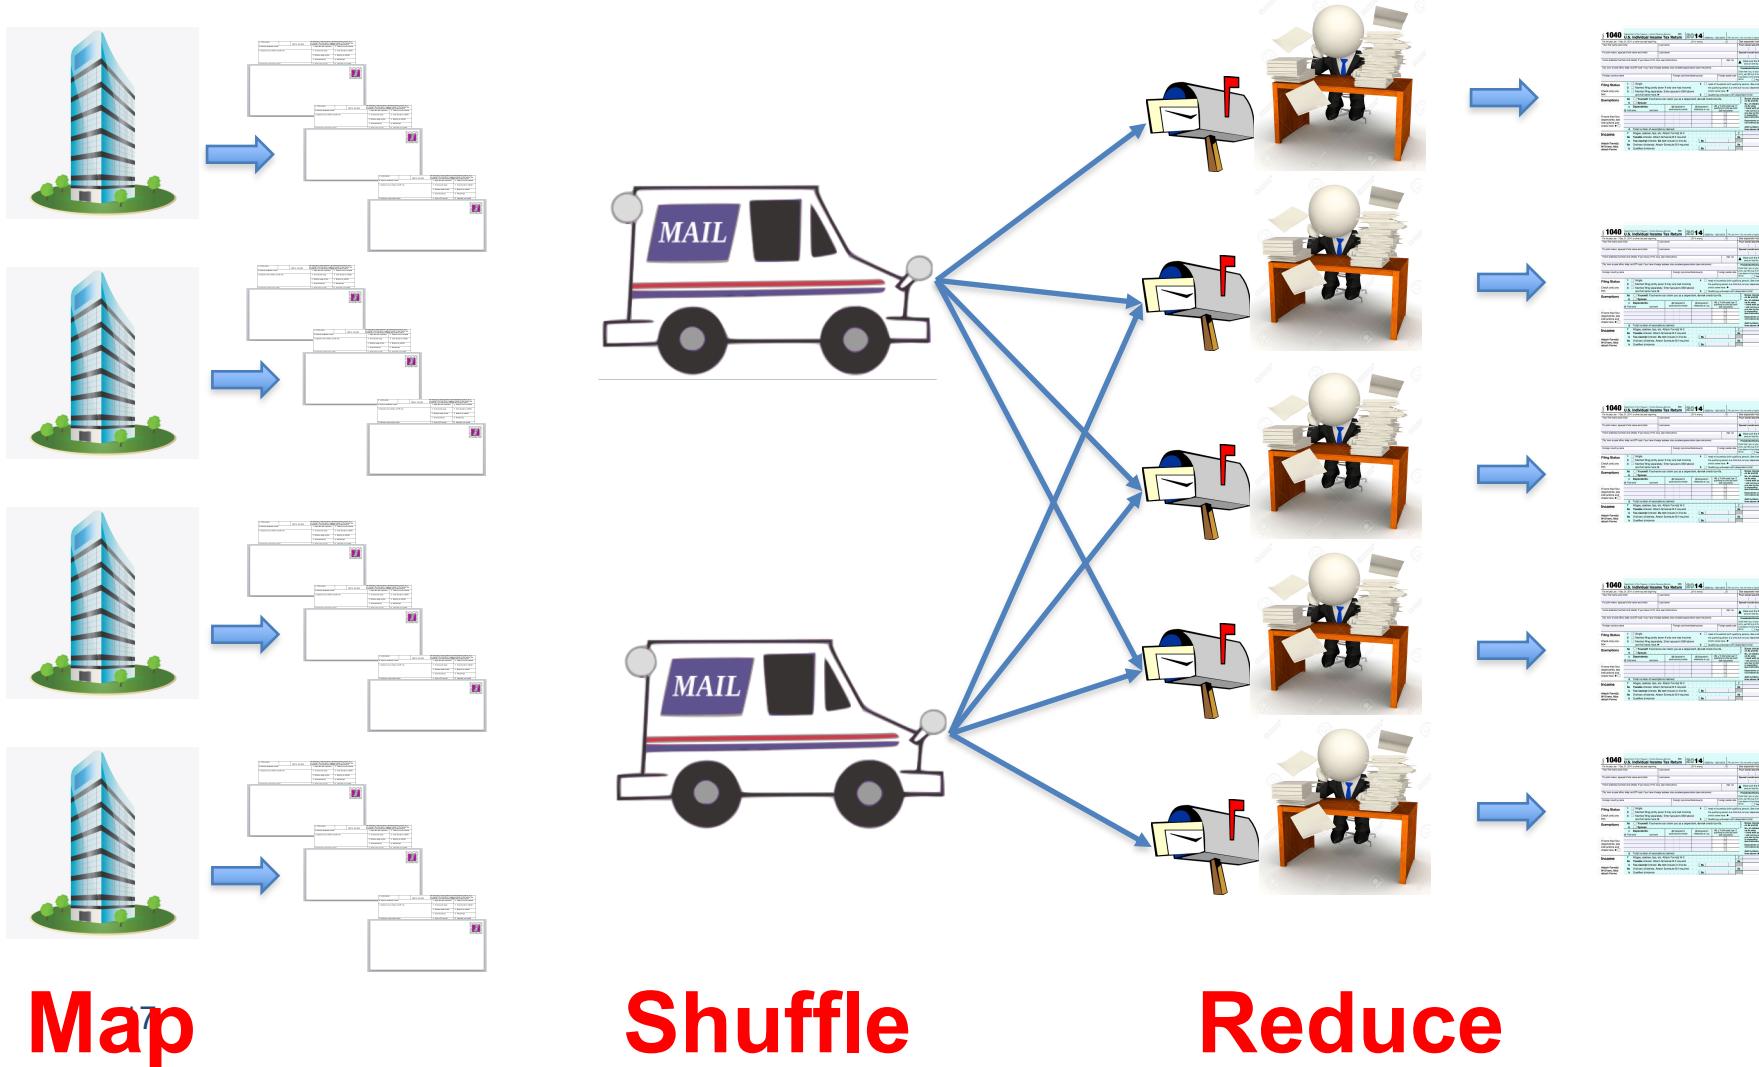
The MapReduce Workflow

Step 3: Reduce

The **reduce** task receives all the **values** with the **same key** and combines them into the final output.



The MapReduce Workflow



Map

- The map function is implemented by the user.
- Each map task processes a single line of an input file at a time.
- Map tasks do not communicate with other map tasks.
- Map tasks communicate with reduce tasks **only** through the content of the value in the KV pair.
- A single map task may emit **any number** of KV pairs (including none).

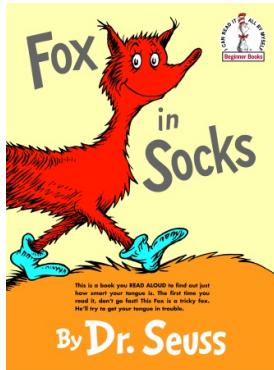
Sort/Shuffle

- The sorting and shuffling process is implemented at the framework level and is opaque to the user.

Reduce

- The reduce function is implemented by the user.
- Each reduce task receives **all the values** associated with a given **key**.
- Reduce tasks do not communicate with each other
- A single reduce task may emit **any number** of result values for each **key** it processes.

WordCount: an example of MapReduce



Map

Key Value
↓ ↓

When tweetle beetles fight, → (When, 1), (**tweetle**, **1**), (beetles, 1), (fight, 1)

it's called a tweetle beetle battle. → (it's, 1), (called, 1), (a, 1), (tweetle, 1), (beetle, 1), (battle, 1)

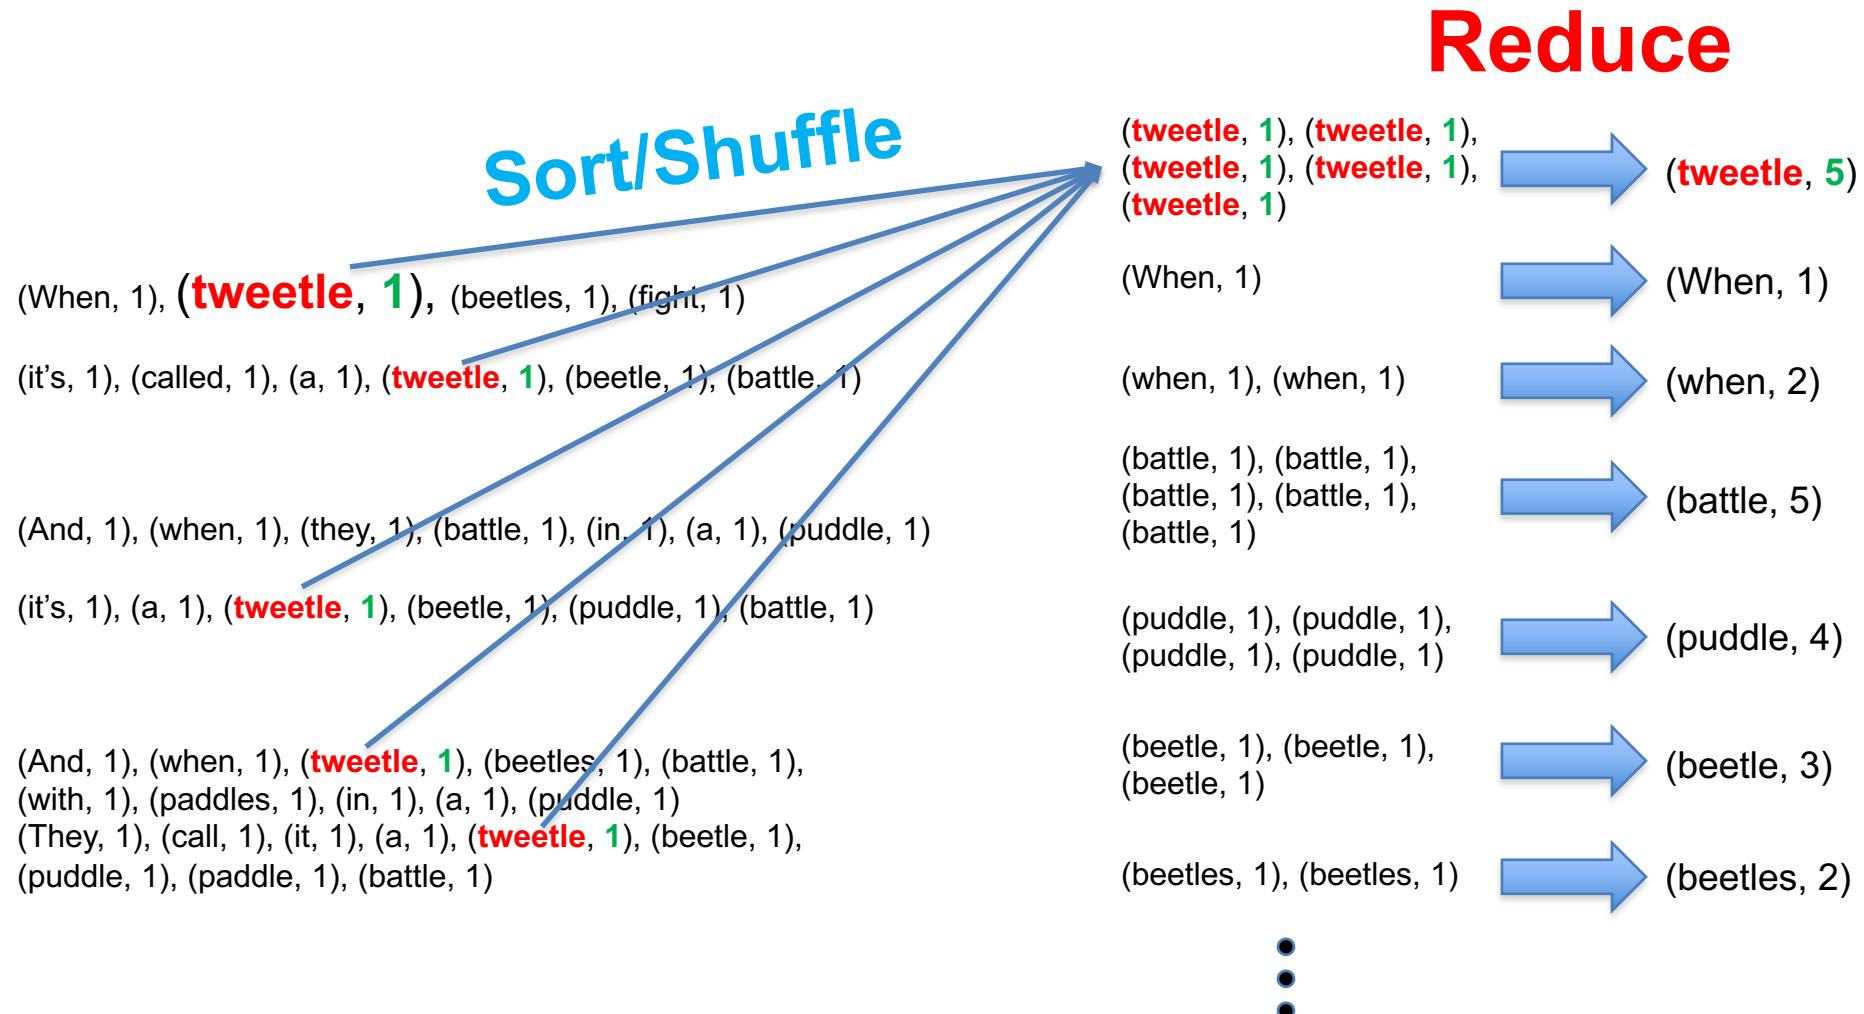
And when they battle in a puddle, → (And, 1), (when, 1), (they, 1), (battle, 1), (in, 1), (a, 1), (puddle, 1)

it's a tweetle beetle puddle battle. → (it's, 1), (a, 1), (tweetle, 1), (beetle, 1), (puddle, 1), (battle, 1)

And when tweetle beetles battle with paddles in a puddle, → (And, 1), (when, 1), (tweetle, 1), (beetles, 1), (battle, 1),
(with, 1), (paddles, 1), (in, 1), (a, 1), (puddle, 1)

They call it a tweetle beetle puddle paddle battle. → (They, 1), (call, 1), (it, 1), (a, 1), (tweetle, 1), (beetle, 1),
(puddle, 1), (paddle, 1), (battle, 1)

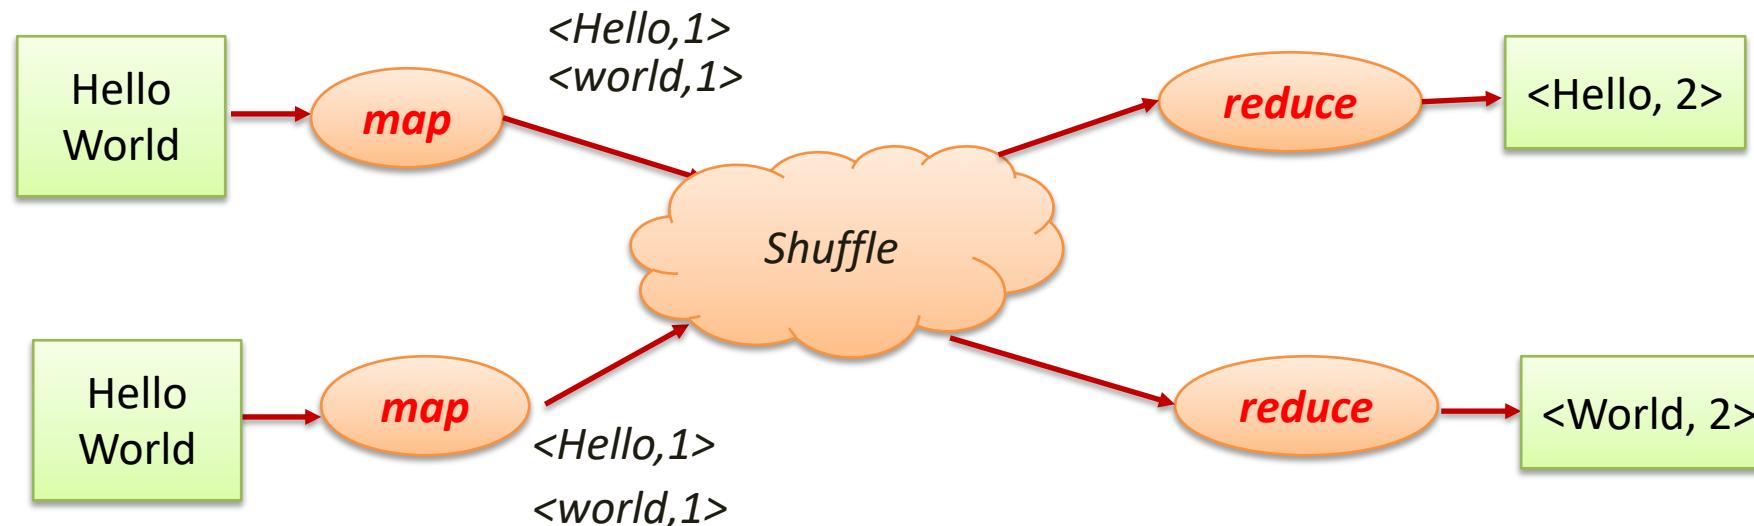
WordCount: an example of MapReduce



Canonical MapReduce

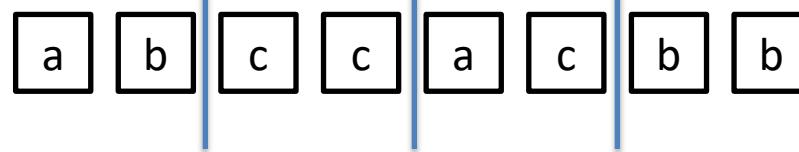
- MapReduce runtime handles the parallel job execution, communication, and data movement
- Users provide *map* and *reduce* functions

Wordcount example:



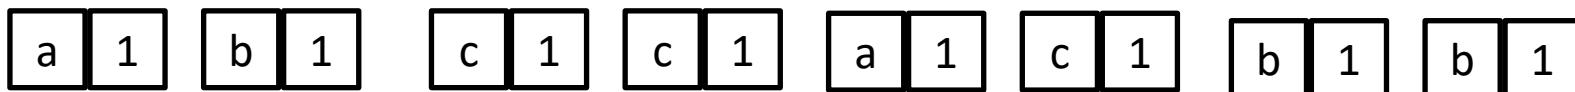
The Canonical MapReduce

Mappers: applied to all
input
data



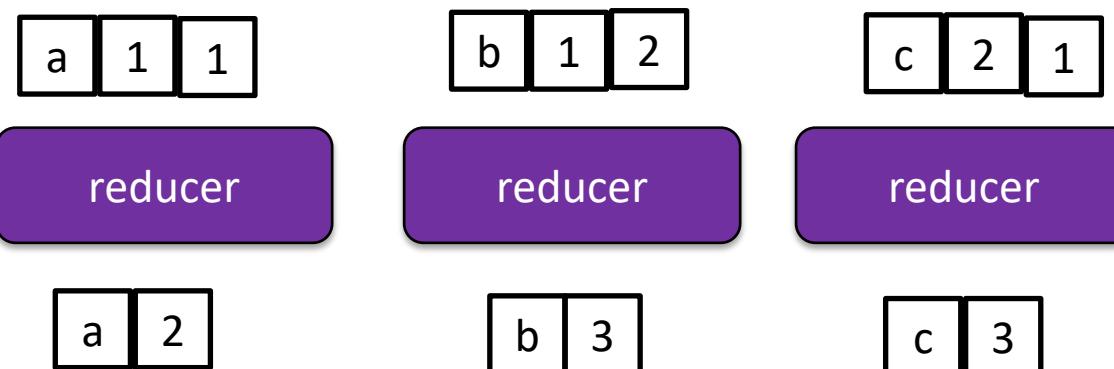
Arbitrary number of
key-value pairs

Barrier:
distributed
sort and group by key



Shuffle and sort: aggregate values by key

Reducers: applied to
all values associated
with the same key



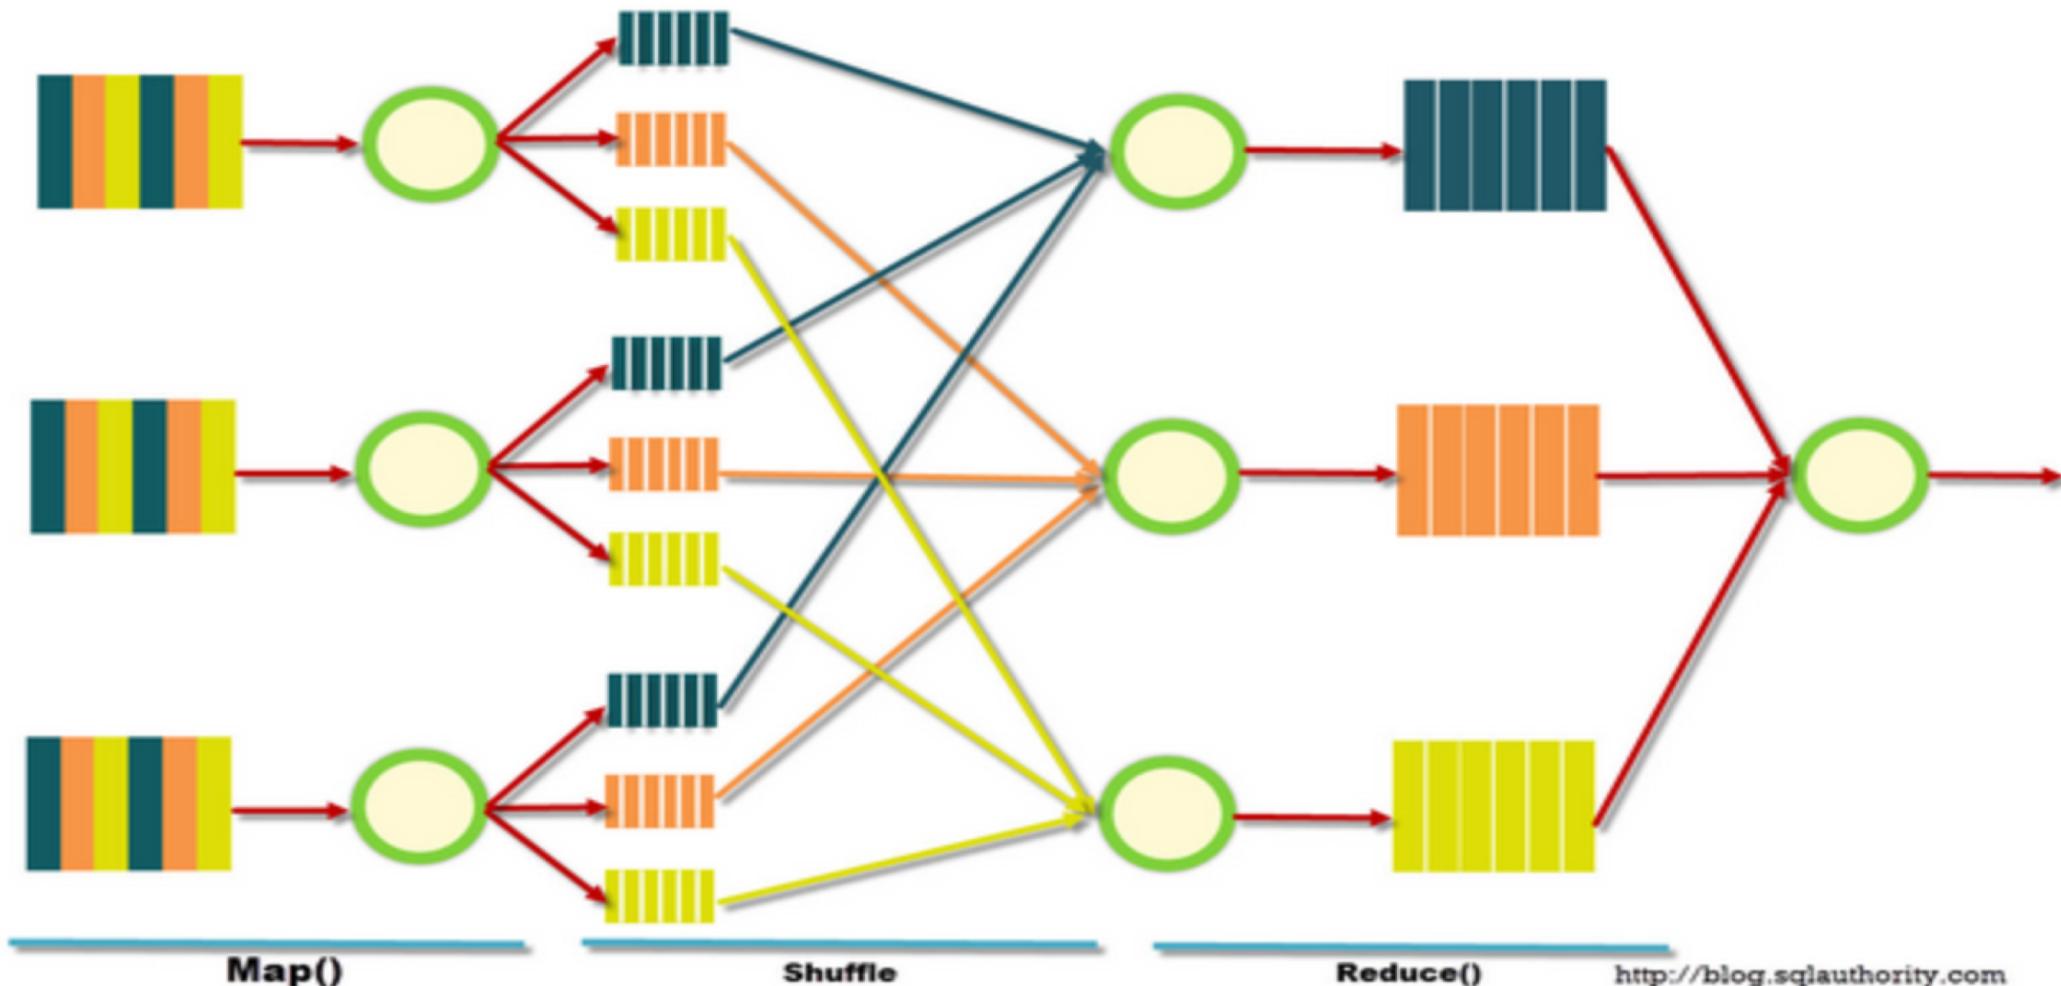
Pseudo-code: WordCount in MapReduce

```
1: class MAPPER
2:     method MAP(docid a, doc d)
3:         for all term t ∈ doc d do
4:             EMIT(term t, count 1)

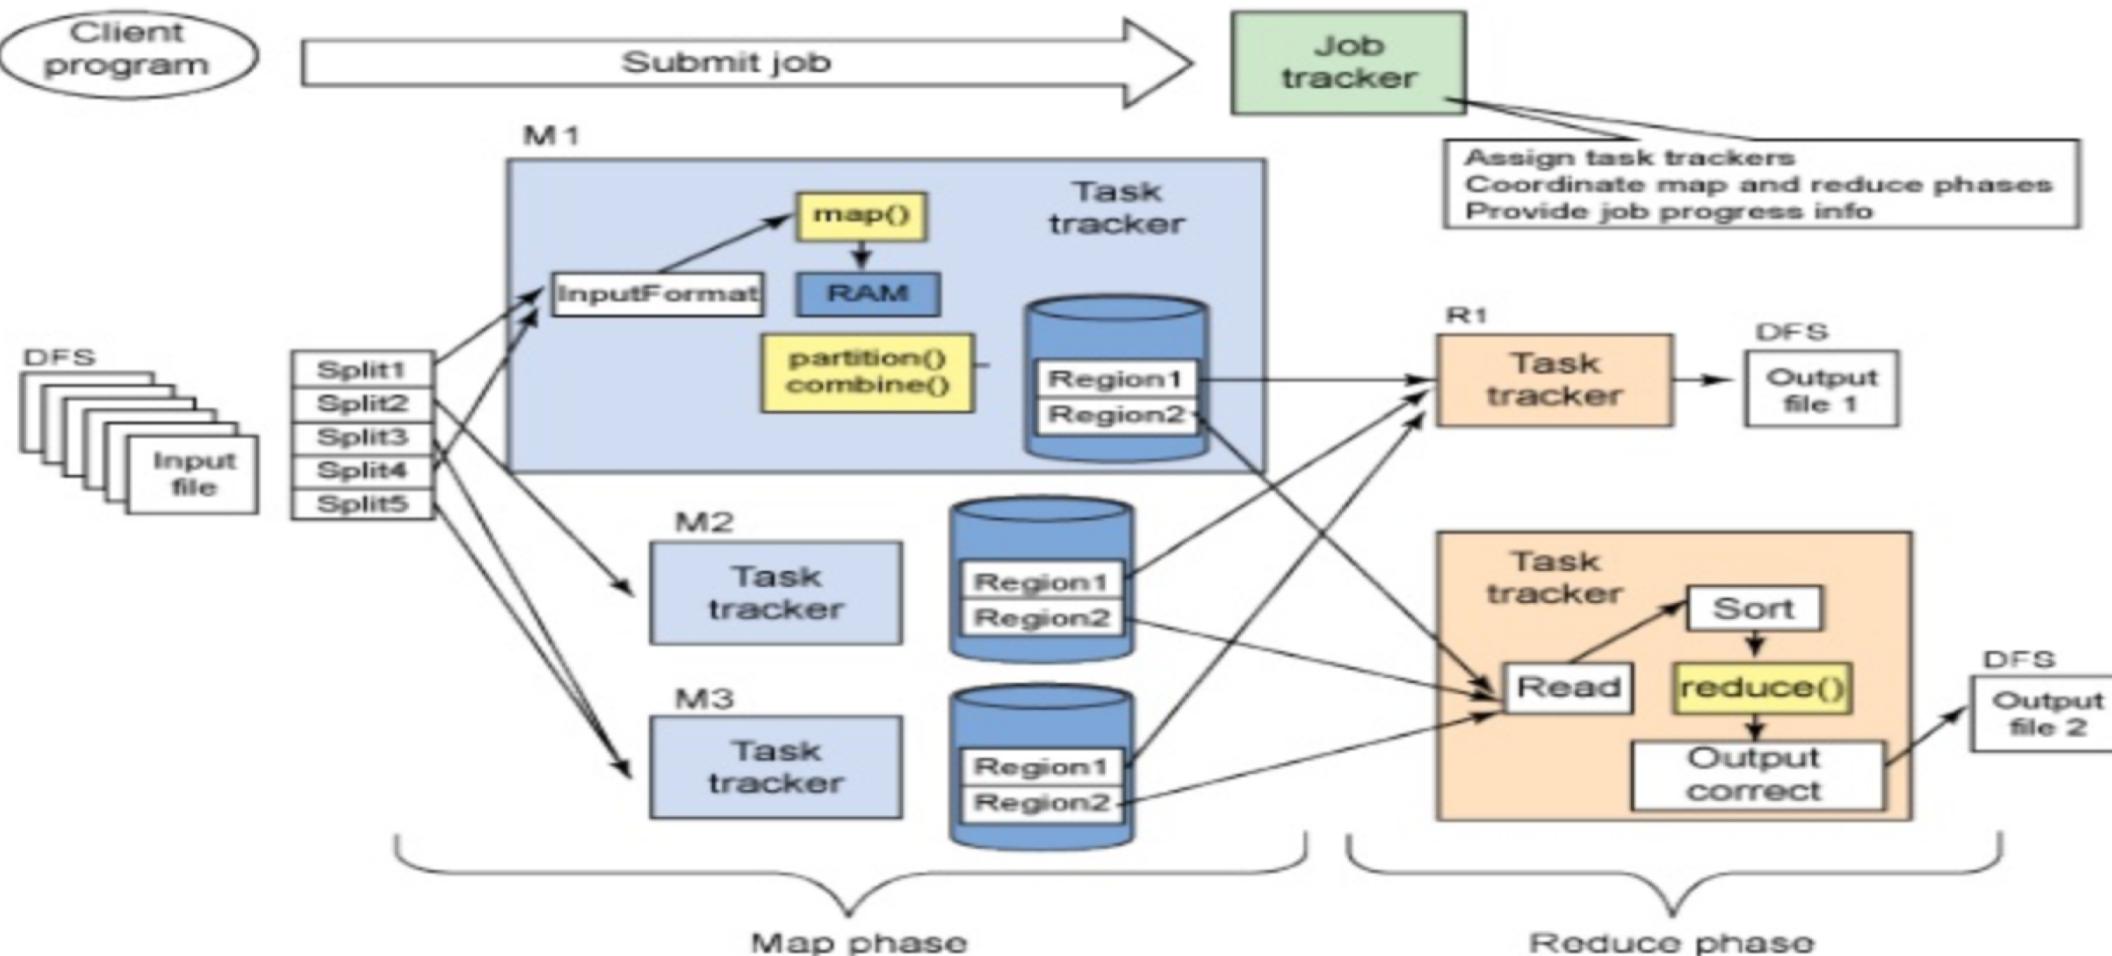
1: class REDUCER
2:     method REDUCE(term t, counts [c1, c2, ...])
3:         sum ← 0
4:         for all count c ∈ counts [c1, c2, ...] do
5:             sum ← sum + c
6:         EMIT(term t, count sum)
```



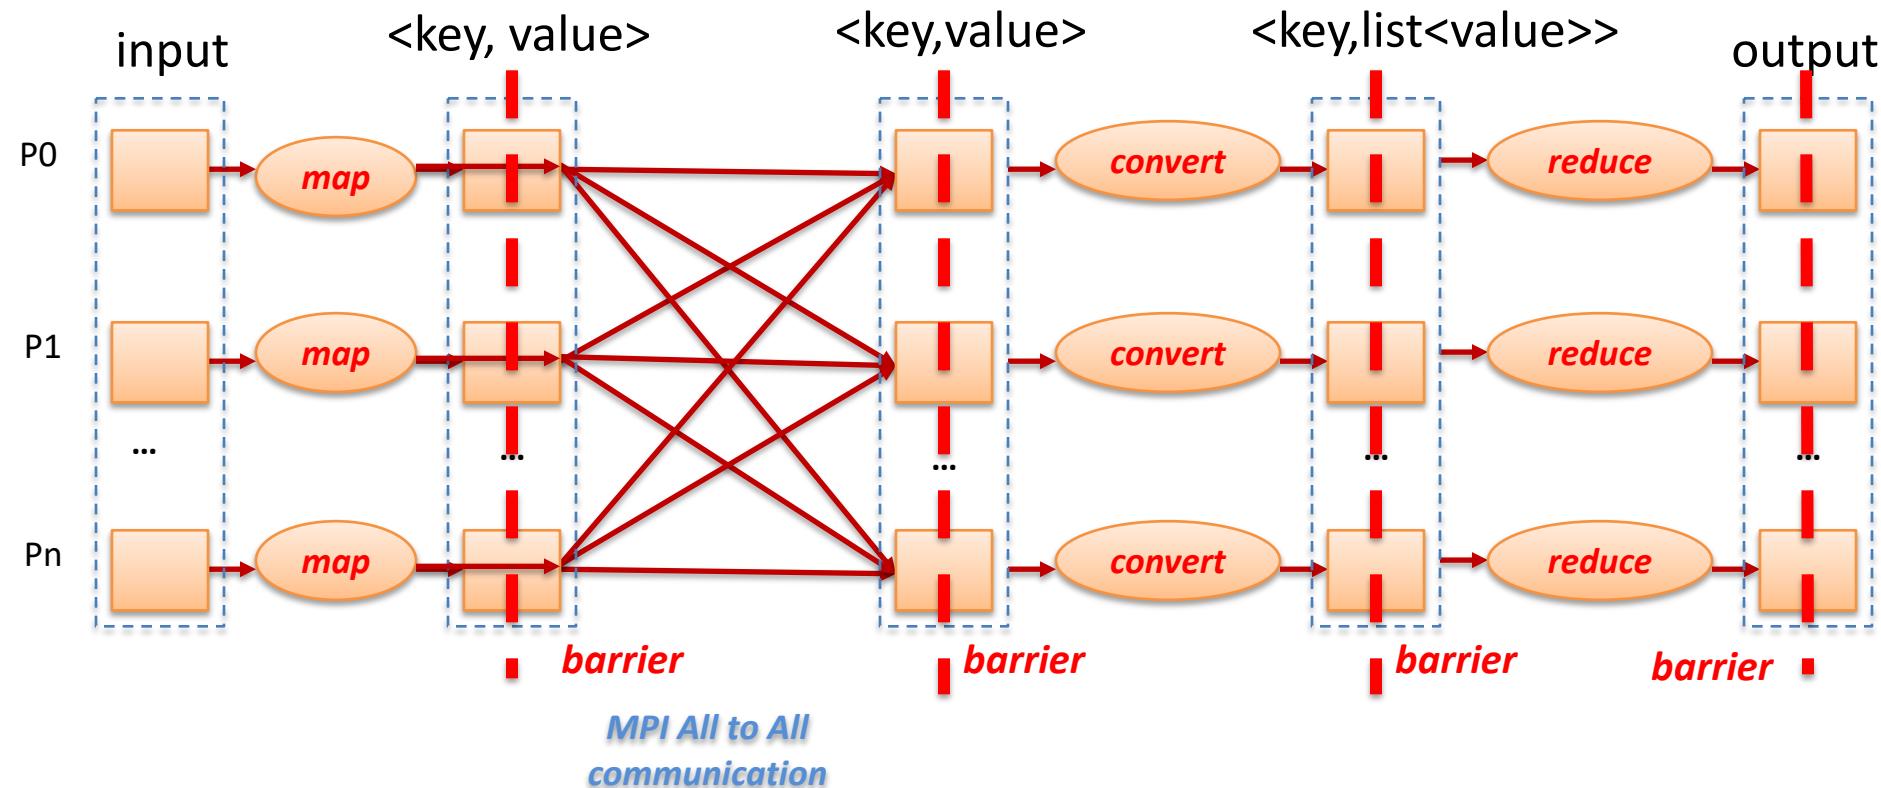
How MapReduce works



Architecture: Hadoop

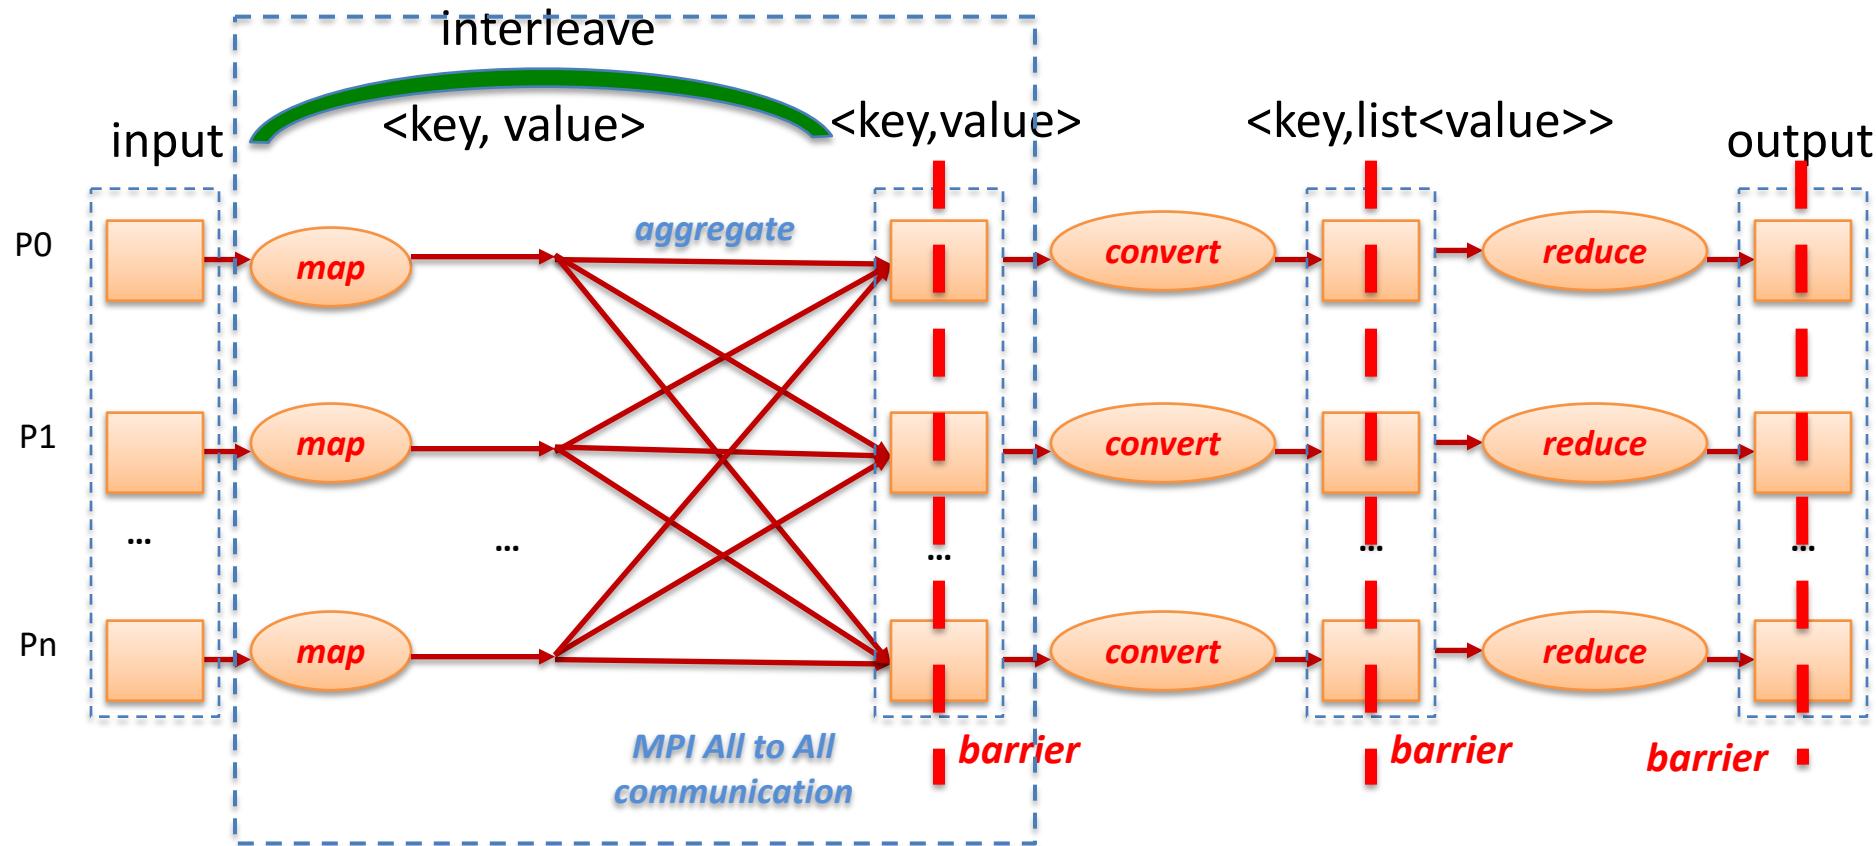


Architecture: MapReduce-MPI



MapReduce in MPI for Large-Scale Graph Algorithms, S. J. Plimpton and K. D. Devine,
Parallel Computing, 37, 610-632 (2011)

Architecture: Mimir



Mimir: Memory-Efficient and Scalable MapReduce for Large Supercomputing Systems.
T. Gao, Y. Guo, B. Zhang, P. Cicotti, Y. Lu, P. Balaji, and M. Taufer. *IPDPS*, 2017

Building Expertise

- Assignment 2



Assignment 2

- Consolidate expertise in Github, Python, and Jupyter Notebook
- Learn to strategize on problems' solutions
 - Write a simple solution first
 - Use git commit to save significant steps
 - Use git branch for new solution efforts
 - Take performance (executing times) into account as the data analyzed grows
 - Learn that problems come with constraints and assumptions

Assignment 2 – Problem 1

(From [ProjectEuler Problem 1](#))

If we list the natural numbers below 10 that are multiples of 3 or 5, we get: 3, 5, 6, and 9. The sum of these multiples is 23.

Write a function that finds the sum of the multiples of p or q below N.

Assumptions and Constraints (I)

- Assume that $1 \leq p < q < N$ and that each of these values are integers. Your code should be able to
 - handle values of N up to at least 100,000,000 (larger and larger data)
 - terminate in less than 1 second (constant execution time!)

Assumptions and Constraints (I)

- Things to keep in mind
 - There are two general approaches to this problem, the naïve (slower) approach and a more mathematical (faster) approach involving the [inclusion-exclusion principle](#). To meet the performance constraints, you will have to implement the fast approach.
 - There are different approaches to measure execution times (wall-clock times). The approach below is one of them. Use what you prefer, as long as you measure wall clock time.



Assignment 2 – Problem 2

- Manipulate a list of names
 - Use the *csv module* to import a list of names
 - Score names in a list based on name “worth” and alphabetical order

Assignment 2 – Problem 2

(From [ProjectEuler Problem 22.](#))

- The file p022_names.txt contains one line with over 5000 comma-separated names, each in all-capital letters and enclosed in quotes.
- Import the names and sort them into alphabetical order.
- Working out the alphabetical value for each name, multiply this value by its alphabetical position in the list to obtain a name score.
- Example:
 - When the list is sorted into alphabetical order, COLIN, which is worth $3 + 15 + 12 + 9 + 14 = 53$, is the 938th name in the list.
 - COLIN would obtain a score of $938 * 53 = 49714$.
- What is the total of the scores for all the names in the file?



Assignment 2 – Problem 3

- Find the smallest TPH number bigger than n
 - Triangular (T), Pentagonal (P), and Hexagonal (H) numbers are generated by given formulae
 - Write the code to find the next triangle number that is also pentagonal and hexagonal

Assignment 2 – Problem 3

- (From [ProjectEuler Problem 45.](#))
- Triangular, Pentagonal, and Hexagonal numbers are generated by the following formulae:

Polygonal	<i>formula for nth term</i>	<i>sequence of terms</i>
Triangular	$T_n = \frac{n(n+1)}{2}$	1, 3, 6, 10, 15, ...
Pentagonal	$P_n = \frac{n(3n-1)}{2}$	1, 5, 12, 22, 35, ...
Hexagonal	$H_n = (2n - 1) n$	1, 6, 15, 28, 45, ...

Assignment 2 – Problem 3

- The number 1 is triangular, pentagonal, and hexagonal (TPH).
- Less obviously, $40755 = T_{285} = P_{165} = H_{143}$ is also TPH
 - 40755 is the smallest TPH number bigger than 1.
- Write a function to find the smallest TPH number bigger than n .
- Use your function to find the smallest TPH bigger than 40755.

Things to consider

- Your choice of data structure can have a significant impact on runtime.
- Think about which operations you are doing the most and choose a data structure which minimizes the average time for this particular operation.
- Python has many built-in data structures
 - the most common data structures are lists, dictionaries, and sets, but Python also has heaps and queues.

Assignment 2 – Problem 4

From [ProjectEuler Problem 87.](#))

- The smallest number expressible as the sum of a prime square, a prime cube, and a prime fourth power is $2^8=2^2+2^3+2^4$.
- There are exactly four numbers below 50 that can be expressed in such a way:

$$2^8 = 2^2 + 2^3 + 2^4$$

$$3^3 = 3^2 + 2^3 + 2^4$$

$$4^9 = 5^2 + 2^3 + 2^4$$

$$4^7 = 2^2 + 3^3 + 2^4$$

Assignment 2 – Problem 4

- Write code to determine the number of positive integers smaller than N that can be written as the sum of a prime square, a prime cube, and a prime fourth power.
- Your code must accept a single command line parameter
 - this time your Jupyter notebook accepts a user input N
- Your code must print a single integer
 - Example: given the input equal to 50, the output is 4

Assumptions and constraints

- For testing, you may assume that N is a positive integer and that $N \leq 50,000,000$ (or larger).
- You should be able to compute the answer when $N = 50,000,000$ and terminate in approximately 1 minutes.

Things to consider

- If you are having a hard time getting started, then break down the problem into smaller manageable pieces.
- Almost certainly you'll need to have a list of primes handy.
 - Can you generate a list of primes?
 - How big do your prime numbers need to be?

For the next week

- Get your solutions for Assignment 2 done before our next class
 - Push your solution into your own repos
 - Suggested reading
 - S. Keshav **How to Read a Paper**
 - Jeffrey Dean and Sanjay Ghemawat (2004) **MapReduce: Simplified Data Processing on Large Clusters**
- Next week
 - More about MapReduce and its optimizations
 - More practical problems

